

```
In [1]: #Name : Ankita Gulde
```

```
#Roll no : 44
```

```
#Section : 3A
```

```
In [1]: #Aim : To perform operation on KNN (K Nearest Neighbor)
```

Importing Libraries

```
In [7]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
from sklearn.model_selection import train_test_split
import warnings
warnings.filterwarnings('ignore')
```

```
In [8]: import os
```

```
In [9]: os.getcwd()
```

```
Out[9]: 'C:\\Users\\HP'
```

```
In [10]: os.chdir("C:\\Users\\HP\\Desktop")
```

```
In [11]: df=pd.read_csv("framingham.csv")
```

```
In [12]: #The "Framingham" heart disease dataset includes over 4,240 records, 15 attrib  
#The goal of the dataset is to predict whether the patient has 10-year risk of
```

```
In [13]: df.head()
```

```
Out[13]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHypertension
0	1	39	4.0	0	0.0	0.0	0	0
1	0	46	2.0	0	0.0	0.0	0	0
2	1	48	1.0	1	20.0	0.0	0	0
3	0	61	3.0	1	30.0	0.0	0	1
4	0	46	3.0	1	23.0	0.0	0	0

```
In [14]: df.describe()
```

```
Out[14]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke
count	4238.000000	4238.000000	4133.000000	4238.000000	4209.000000	4185.000000	4238.000000
mean	0.429212	49.584946	1.978950	0.494101	9.003089	0.029630	0.016958
std	0.495022	8.572160	1.019791	0.500024	11.920094	0.169584	0.016958
min	0.000000	32.000000	1.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	42.000000	1.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	49.000000	2.000000	0.000000	0.000000	0.000000	0.000000

In	df.info()						
max	1.000000	70.000000	4.000000	1.000000	70.000000	1.000000	1.0
	<pre><class 'pandas.core.frame.DataFrame' RangeIndex: 4238 entries, 0 to 4237 Data columns (total 16 columns): " " " " " " " " " " " " " " " "</pre> Dtype						
male		age	education	currentSmoker	cigsPerDay	BPMeds	prevalentS
75%	1.000000	56.000000	3.000000	1.000000	20.000000	0.000000	0.0

```
15    TenYearCHD      423 non-    int64
      8             null
```

```
memory usage:      KB
529.9
```

```
df.isna().sum()
```

```
male      0
age       0
education 105
currentSmoker  0
cigsPerDay 29
BPMeds     53
prevalentStroke  0
prevalentHyp  0
diabetes    0
totChol     50
sysBP       0
diaBP       0
BMI         19
heartRate    1
glucose     388
TenYearCHD   0
dtype: int64
dtypes: float64(9), int64(7)
```

```
In [16]:
```

```
Out[16]:
```

```
In [ ]: #Since, only a few rows have null values in them, we are only removing those
#df = df.dropna(subset=['heartRate','BMI','cigsPerDay','totChol','BPMeds'])
```

```
In [17]: df
```

Out [17]:

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp	d
	0	1	39	4.0	0	0.0	0.0	0	0

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHypertension
1	0	46	2.0	0	0.0	0.0	0	0
2	1	48	1.0	1	20.0	0.0	0	0
3	0	61	3.0	1	30.0	0.0	0	1
4	0	46	3.0	1	23.0	0.0	0	0
...
4233	1	50	1.0	1	1.0	0.0	0	1
4234	1	51	3.0	1	43.0	0.0	0	0
4235	0	48	2.0	1	20.0	NaN	0	0
4236	0	44	1.0	1	15.0	0.0	0	0
4237	0	52	2.0	0	0.0	0.0	0	0

4238 rows × 9 columns

Missing Value Treatment

Since, 'glucose' and 'education' columns had a significant amount of null values, so we replaced them with the mean of values for their respective columns

```
In [18]: df['glucose'].fillna(value = df['glucose'].mean(),inplace=True)

In [19]: df['education'].fillna(value = df['education'].mean(),inplace=True)

In [20]: df['heartRate'].fillna(value = df['heartRate'].mean(),inplace=True)

In [21]: df['BMI'].fillna(value = df['BMI'].mean(),inplace=True)

In [22]: df['cigsPerDay'].fillna(value = df['cigsPerDay'].mean(),inplace=True)

In [23]: df['totChol'].fillna(value = df['totChol'].mean(),inplace=True)

In [24]: df['BPMeds'].fillna(value = df['BPMeds'].mean(),inplace=True)

In [26]: df.isna().sum()

Out[26]: male                0
age                0
education          0
currentSmoker      0
cigsPerDay         0
BPMeds             0
prevalentStroke    0
```

prevalent
Hyp
0
diabetes
0
totChol
0
sysBP
0
diaBP
0
BMI
0

```
heartRate      0
glucose        0
TenYearCHD     0
dtype: int64
```

Logistic Regression Model

```
In [27]: #Splitting the dependent and independent variables.
x = df.drop("TenYearCHD",axis=1)
y = df['TenYearCHD']
```

```
In [28]: x #checking the features
```

```
Out[28]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHypertension
0	1	39	4.0	0	0.0	0.00000	0	0
1	0	46	2.0	0	0.0	0.00000	0	0
2	1	48	1.0	1	20.0	0.00000	0	0
3	0	61	3.0	1	30.0	0.00000	0	1
4	0	46	3.0	1	23.0	0.00000	0	0
...
4233	1	50	1.0	1	1.0	0.00000	0	1
4234	1	51	3.0	1	43.0	0.00000	0	0
4235	0	48	2.0	1	20.0	0.02963	0	0
4236	0	44	1.0	1	15.0	0.00000	0	0
4237	0	52	2.0	0	0.0	0.00000	0	0

4238 rows × 15 columns

Train Test Split

```
In [29]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2,random_stat
```

```
In [30]: y_train
```

```
Out[30]: 3252    0
          3946    0
          1261    0
          2536    0
```



```
4089    0
      ..
3444    0
466     0
3092    0
3772    0
860     0
```

```
Name      TenYearCHD, Length: 3390, dtype:
:          int64
```

KNN Clas sifier

```
In [31]: from sklearn.neighbors import KNeighborsClassifier
         knn = KNeighborsClassifier(n_neighbors=5, p=2, metric='minkowski')
         knn.fit(x_train, y_train)
         acc = knn.score(x_test, y_test)*100
         print(acc)
```

83.13679245283019

In []: